

A Mini Project Report

On

**Stock Price prediction using advanced LSTM techniques and
real time forecasting**

*Submitted to CMREC (UGC Autonomous), Affiliated to JNTUH
In Partial Fulfillment of the requirements for the Award of Degree of*

BACHELOR OF TECHNOLOGY

IN

COMPUTER SCIENCE AND ENGINEERING (AI&ML)

Submitted By

G. BHARGAV SAI	(228R1A6625)
K. SAI MANASA	(228R1A6630)
K. SWATHI	(228R1A6633)
M. RISHIKA SHIVANI	(228R1A6637)

Under the Esteemed guidance of

Mr. B. Sai Kumar

Assistant Professor, Department of CSE (AI & ML)



**Department of Computer Science & Engineering
(AI&ML)**

CMR ENGINEERING COLLEGE
UGC AUTONOMOUS

(Accredited by NAAC&NBA, Approved by AICTE, NEW DELHI, Affiliated to
JNTU, Hyderabad, Kandlakoya, Medchal Road, R.R. Dist. Hyderabad-501 401)

2024-2025

CMR ENGINEERING COLLEGE

UGC AUTONOMOUS

(Accredited by NAAC&NBA, Approved by AICTE, NEW DELHI, Affiliated to JNTU, Hyderabad, Kandlakoya, Medchal Road, R.R. Dist. Hyderabad-501 401)

Department of Computer Science & Engineering (AI & ML)



CERTIFICATE

This is to certify that the project entitled “**Stock Price Prediction using advanced LSTM techniques and real time forecasting**” is a bonafide work carried out by

G.BHARGAV	(228R1A6625)
K.SAI MANASA	(228R1A6630)
K.SWATHI	(228R1A6633)
M.RISHIKA	(228R1A6637)

in partial fulfillment of the requirement for the award of the degree of **BACHELOR OF TECHNOLOGY** in **COMPUTER SCIENCE AND ENGINEERING (AI&ML)** from CMR Engineering College, under our guidance and supervision.

The results presented in this project have been verified and are found to be satisfactory. The results embodied in this project have not been submitted to any other university for the award of any other degree or diploma.

Internal Guide Mr.
B. Sai Kumar
Assistant Professor
CSE (AI&ML)

Project Coordinator
Mrs. M. Srikala
Assistant Professor
CSE (AI&ML)

Head of the Department
Dr. Madhavi Pingili
Professor & HOD
CSE (AI&ML)

DECLARATION

This is to certify that the work reported in the present project entitled “**Stock Price Prediction using advanced LSTM techniques and real time forecasting**” is a record of bonafide work done by us in the Department of Information Technology, CMR Engineering College, JNTU Hyderabad. The reports are based on the project work done entirely by us and not copied from any other source. We submit our project for further development by any interested students who share similar interests to improve the project in the future.

The results embodied in this project report have not been submitted to any other University or Institute for the award of any degree or diploma to the best of our knowledge and belief.

G.BHARGAV

(228R1A6625)

K.SAI MANASA

(228R1A6630)

K.SWATHI

(228R1A6633)

M.RISHIKA

(228R1A6637)

ACKNOWLEDGMENT

We are extremely grateful to **Dr. A. Srinivasula Reddy**, Principal and **Dr. Madhavi Pingili**, HOD, **Department of CSE (AI & ML), CMR Engineering College** for their constant support.

We are extremely thankful to **Mr. B. Sai Kumar**, Assistant Professor, Internal Guide, Department of CSE (AI & ML), for his constant guidance, encouragement and moral support throughout the project.

We will be failing in duty if we do not acknowledge with grateful thanks to the authors of the references and other literatures referred in this Project.

We thank **Mrs.M.Srikala**, Mini Project Coordinator for her constant support in carrying out the project activities and reviews.

We express our thanks to all staff members and friends for all the help and co-ordination extended in bringing out this project successfully in time.

Finally, We are very much thankful to our parents who guided us for every step.

G.BHARGAV	(228R1A6625)
K.SAI MANASA	(228R1A6630)
K.SWATHI	(228R1A6633)
M. RISHIKA	(228R1A6637)

CONTENTS

TOPIC	PAGE NO
ABSTRACT	I
LIST OF FIGURES	II
LIST OF TABLES	III
1. INTRODUCTION	4
1.1. Introduction & Objectives	4
1.2. Project Objectives	5
1.3. Purpose of the project	5
1.4. Existing System with Disadvantages	5
1.5. Proposed System With features	6
1.6. Input and Output Design	8
2. LITERATURE SURVEY	10
3. SOFTWARE REQUIREMENT ANALYSIS	14
3.1. Problem Specification	14
3.2. Modules and their Functionalities	14
3.3. Non-Functional Requirements	15
34. Feasibility Study	16
4. SOFTWARE & HARDWARE REQUIREMENTS	17
4.1. Software requirements	17
4.2. Hardware requirements	17

5. SOFTWARE DESIGN	18
5.1. System Architecture	18
5.2. Dataflow Diagrams	18
5.3. UML Diagrams	19
6. CODING AND IMPLEMENTATION	25
6.1. Source code	25
6.2. Implementation	45
6.2.1. Python	45
6.2.2 Machine Learning	46
7. SYSTEM TESTING	47
7.1. Types of System Testing	47
7.2. Test Cases	50
8. OUTPUT SCREENS	51
9. CONCLUSION	55
10. FUTURE ENHANCEMENTS	56
11. REFERENCES	57

ABSTRACT

The dynamic and volatile nature of financial markets poses significant challenges for accurate stock price forecasting. Traditional statistical models often struggle to capture the complex, non-linear dependencies inherent in time series stock data. This study explores a scalable and robust approach to time series forecasting of stock prices using advanced Long Short-Term Memory (LSTM) techniques, which have proven effective in modeling sequential data with temporal dependencies. We propose an enhanced LSTM-based framework that integrates attention mechanisms, multi-layered architectures, and feature engineering to improve prediction accuracy and model generalization across various stock indices. Our model incorporates technical indicators such as moving averages, RSI, and MACD, along with historical stock prices and trading volumes, to enrich the input feature space. We employ a sliding window strategy for time series segmentation, allowing the model to learn patterns over fixed intervals. To ensure scalability, the architecture is optimized to handle large volumes of high-frequency data using parallel training on distributed computing platform.

Experimental evaluations were conducted on benchmark datasets from major stock exchanges, including NASDAQ and NYSE. The proposed LSTM model outperformed traditional methods such as ARIMA and vanilla LSTM, demonstrating lower mean squared error (MSE) and higher directional accuracy. Furthermore, the attention-enhanced LSTM layers enabled the model to focus on relevant time steps, thereby improving interpretability and decision-making support. This research highlights the potential of advanced LSTM techniques in building scalable and accurate stock forecasting systems. The proposed methodology can be extended to other financial time series applications, offering valuable insights for traders, investors, and financial analysts seeking data-driven forecasting solutions in increasingly complex markets.

LIST OF FIGURES

S.NO	FIGURE NO	DESCRIPTION	PAGENO
1	1.5.1	Block diagram of proposed system	5
2	5.1	System Architecture	16
3	5.2	Data Flow diagram	17
4	5.3.1	Sequence diagram	19
5	5.3.2	Use case diagram	20
6	5.3.3	Activity diagram	21
7	5.3.4	Class diagram	22
8	8.1	Output Screen-1	49
9	8.2	Output Screen-2	50
10	8.3	Output Screen-3	50
11	8.4	Output Screen-4	51
12	8.5	Output Screen-5	51
13	8.6	Output Screen-6	52

LIST OF TABLES

S.NO	TABLE NO	DESCRIPTION	PAGENO
1	7.2	Test Cases	49

1. INTRODUCTION

1.1 Introduction & Objectives

Forecasting stock prices is a complex and high-stakes task in the financial domain, where accurate predictions can significantly influence investment decisions. The stock market is inherently volatile, driven by numerous dynamic and non-linear factors such as investor sentiment, economic indicators, and global events. Traditional statistical models like ARIMA, though useful in linear trend forecasting, often fail to capture these complexities and long-term dependencies in time series data. With the advent of deep learning, Long Short-Term Memory (LSTM) networks have emerged as a powerful tool for modeling sequential data. LSTMs are capable of learning temporal patterns and dependencies over long time horizons, making them particularly suitable for stock market forecasting.

However, standard LSTM models face challenges in scalability and interpretability, especially when applied to large datasets with high-frequency trading information. Moreover, they often do not fully utilize the wealth of available features such as technical indicators and volume trends. This study proposes a scalable and advanced LSTM-based framework for stock price forecasting. By incorporating attention mechanisms, multi-layer architectures, and enriched feature sets, the model aims to improve prediction accuracy while remaining computationally efficient and interpretable.

The first objective is to design a scalable LSTM-based model capable of processing large volumes of time series stock data efficiently. Secondly, the model aims to enhance forecasting accuracy through the integration of deep learning enhancements like attention mechanisms and multi-layer LSTM networks. A third objective is to leverage additional features such as technical indicators and trading volumes to provide richer input representations. The study also seeks to benchmark the proposed approach against traditional models like ARIMA and basic LSTM to evaluate performance improvements. Lastly, it aims to provide interpretable outputs that highlight key influencing factors in predictions, aiding traders and analysts in making more informed, data-driven decisions.

1.2 Project Objectives

The primary objective of this project is to develop an advanced stock price prediction system using deep learning techniques, specifically Long Short-Term Memory (LSTM) models. The system aims to analyze historical stock data and predict future prices with improved accuracy by incorporating Bidirectional, Stacked, and Attention-based LSTM architectures. Additionally, the project seeks to integrate real-time data fetching and deliver predictions through an interactive web-based dashboard, enabling users to make more informed investment decisions.

1.3 Purpose of the Project

The purpose of this project is to design and implement a reliable stock price prediction system that leverages deep learning models to forecast future stock trends. By utilizing historical stock market data and advanced LSTM-based architectures, the project aims to assist investors, analysts, and financial enthusiasts in making data-driven decisions. The system also incorporates real-time data updates and a user-friendly interface to ensure accessibility, timely insights, and practical application in real-world trading and investment scenarios.

1.4 Existing System with Disadvantages

Traditional stock price prediction systems often rely on statistical methods such as Moving Averages, ARIMA models, or basic machine learning algorithms like Linear Regression and Random Forests. While these methods can capture general trends, they struggle to model the highly non-linear and volatile nature of financial time series data. Moreover, many existing systems depend heavily on manually engineered features and fail to consider temporal dependencies effectively. They also lack real-time data integration and often provide static, delayed predictions. As a result, these systems offer limited accuracy and adaptability, making them less effective for dynamic market conditions.

1.5 Proposed System with Features

The proposed system aims to build a scalable, accurate, and intelligent time series forecasting model for predicting stock prices using advanced LSTM (Long Short-Term Memory) architectures. It is designed to address the shortcomings of existing systems by incorporating deep learning techniques capable of capturing the non-linear, complex, and temporal nature of financial time series data. The system is not only focused on improving prediction accuracy but also emphasizes scalability, real-time adaptability, and integration with external market factors.

At the core of the system is an enhanced LSTM architecture. Unlike traditional unidirectional LSTMs, this system employs Bidirectional LSTM (BiLSTM) layers, which process input sequences in both forward and backward directions. This helps the model understand future and past context more effectively, improving the learning of sequential patterns.

To enhance forecasting performance across multiple time horizons and stock tickers, the system adopts a Sequence-to-Sequence (Seq2Seq) framework, which is well-suited for modeling variable-length input and output sequences. This allows the model to generate multi-step forecasts, offering more flexibility and foresight in prediction tasks. The model is also designed to accept multi-feature inputs, such as historical stock prices, trading volumes, technical indicators (e.g., moving averages, RSI), and macroeconomic indicators, thus broadening its predictive capacity.

A key innovation in the proposed system is its integration with external data sources. Market sentiment data derived from news articles, social media feeds, and financial reports are preprocessed using natural language processing (NLP) techniques and fed into the model alongside traditional numerical inputs. This multi-modal approach helps the model understand not only numerical trends but also investor sentiment and external market forces, which often have a significant impact on stock price movements.

To ensure scalability, the system is built using distributed computing frameworks and is optimized for deployment on cloud-based platforms. It can handle large volumes of data across hundreds of stocks and is designed for real-time prediction and continuous learning. This enables the system to update itself as new data becomes available, adapting quickly to market changes and improving over time.


STOCK PRICE PREDICTION

Stock price prediction is a significant challenge in the domain of financial analytics and artificial intelligence. The goal is to forecast future stock prices based on historical data, patterns, and sometimes external factors like news sentiment or macroeconomic indicators. Accurate stock prediction models can be highly beneficial for investors, traders, and financial institutions to make informed decisions and minimize risks.

Stock prices are influenced by a complex set of variables including company performance, investor sentiment, market trends, geopolitical events, and economic indicators. Due to the highly volatile and non-linear nature of financial markets, traditional models often fall short in capturing intricate patterns and dependencies.

LSTM models are designed to remember long-term dependencies, making them highly suitable for modeling sequential data like stock prices. Enhancements such as Bidirectional LSTMs, which process data in both forward and backward directions, Stacked LSTMs for deeper feature learning, and Attention Mechanisms for focusing on key time steps have further improved prediction performance.

This project leverages these advanced LSTM-based architectures along with real-time stock data collection to build a dynamic and interactive prediction system. By integrating the model into a web interface, users can view forecasts and trends conveniently, making this project both technically robust and practically useful.



Advantages

- Predicts future stock prices accurately using advanced LSTM models.
- Works well with time-series data like stock prices.
- Can update predictions with real-time data.
- Cleans and prepares data automatically for better results.
- Shows how good the model is using clear error measurements.
- Displays predictions in easy-to-understand graphs.
- Available as a web app that anyone can access.
- Helps investors make smarter decisions based on data.

1.6 Input and Output Design

1.7 Design Input Design

The input design defines the structure, format, and source of data used to train and evaluate the stock price prediction model. In this project, historical stock market data is collected using the yfinance API, which provides features such as opening price, closing price, highest and lowest prices, trading volume, and date. The input data undergoes preprocessing steps including handling missing values, normalization using MinMaxScaler, and sequence generation for LSTM model compatibility. The input is structured into fixed-size time windows to capture temporal patterns and dependencies. Additionally, real-time input data is fetched periodically for live predictions, ensuring that the model can respond to current market trends. This well-structured input design ensures that the deep learning model receives clean, consistent, and meaningful data for accurate forecasting.

Objectives

- Develop a stock price prediction model using advanced LSTM architectures (Bidirectional, Stacked, Attention-based).
- Analyze and preprocess historical stock market data for effective learning.

- Incorporate real-time stock data fetching to enable up-to-date forecasting.
- Evaluate model performance using metrics like RMSE and MAE to ensure accuracy.
- Build an interactive web application for users to visualize and access predictions easily.
- Help investors and traders make informed decisions based on reliable price forecasts.

Output Design

The output design specifies how the predicted stock price information and related results are formatted, communicated, and utilized by the end-users or other systems. The primary output of this project is the predicted stock price values for future time points, generated by the deep learning LSTM-based model. These predicted values are organized as a sequential time series that corresponds to future dates beyond the training data.

In addition to raw predictions, the system outputs key evaluation metrics such as Root Mean Squared Error (RMSE) and Mean Absolute Error (MAE). These metrics provide quantitative measures of the model's prediction accuracy and are essential for validating and fine-tuning the model.

The outputs are presented visually through dynamic and interactive charts, typically implemented using visualization libraries like Matplotlib, Seaborn, or Plotly. These visualizations compare historical actual prices against the predicted prices, allowing users to easily interpret trends, forecast potential market movements, and identify periods of volatility.

Moreover, the output system supports real-time updates by fetching fresh market data continuously and re-running predictions as needed. This ensures that the displayed forecasts remain relevant and actionable.

Finally, the prediction results and visualizations are integrated into a user-friendly web application interface, built with frameworks such as Flask or Streamlit. This interface provides stakeholders, investors, or analysts with convenient access to the forecast data, enabling them to make timely and informed decisions in stock trading or portfolio management.

2. LITERATURE SURVEY

- [1] **Box, G.E.P., Jenkins, G.M., Reinsel, G.C. (2019).** *Time Series Analysis: Forecasting and Control*. Box and colleagues introduced traditional statistical models like ARIMA for time series forecasting. These models work well for stationary and linear data, offering a solid mathematical foundation. However, stock prices are often non-linear and volatile, making these models less effective. Additionally, ARIMA requires rigorous preprocessing for stationarity, which is unrealistic for real-world financial data.
- [2] **Kim, H.Y. (2019).** *Financial Time Series Forecasting Using Support Vector Machines. Neurocomputing*, 55(1-2), 307-319. To overcome the limitations of statistical methods, Kim proposed the use of Support Vector Machines (SVM) for stock prediction. These machine learning models handle non-linearity better and use technical indicators as features. However, they treat each data point independently, failing to capture time-based dependencies which can lead to further process.
- [3] **Fischer, T., Krauss, C. (2018).** *Deep Learning With Long Short-Term Memory Networks for Financial Market Predictions. European Journal of Operational Research*, 270(2), 654-669. Fischer and Krauss applied LSTM networks to stock price prediction, which could learn temporal sequences and long-range dependencies. LSTMs overcome the vanishing gradient problem and reduce reliance on manual features. This makes them effective for modeling noisy, complex financial data.
- [4] **Fischer, T., Krauss, C. (2018).** *Deep Learning With Long Short-Term Memory Networks for Financial Market Predictions*. Their extended study compared LSTM performance with ARIMA and Random Forests on S&P 500 stocks. LSTMs achieved higher accuracy by better capturing stock price dynamics. However, the study lacked exogenous data like news or economic indicators. It also did not address scalability or real-time deployment for live trading.
- [5] **Qin, Y., Song, D., Chen, H., Cheng, W., Jiang, G., Cottrell, G.W. (2018).** *A Dual-Stage Attention-Based Recurrent Neural Network for Time Series Prediction*. Qin and colleagues improved LSTM models by adding an attention mechanism. This allows the model to focus on the most relevant time steps during forecasting. It improves accuracy by filtering out noise and capturing key temporarily.

- [6] **Schuster, M., Paliwal, K.K. (2017). *Bidirectional Recurrent Neural Networks. IEEE Transactions on Signal Processing*, 45(11), 2673-2681.** They introduced Bidirectional RNNs, which process input in both forward and backward directions. BiLSTMs improve prediction by using both past and future context. This bidirectional flow captures more complete temporal patterns in the data. Such models are particularly useful in complex stock market environments.
- [7] **Sutskever, I., Vinyals, O., Le, Q.V. (2017). *Sequence to Sequence Learning with Neural Networks*.** This work proposed Sequence-to-Sequence (Seq2Seq) models for sequence prediction tasks. The encoder processes historical data, and the decoder forecasts multiple future values. This enables multi-horizon stock price prediction. Seq2Seq models are flexible and effective for forecasting over days or weeks.
- [8] **Zhang, L., Aggarwal, C.C., Qi, G.J. (2017). *Stock Price Prediction via Discovering Multi-Frequency Trading Patterns*.** Zhang and team developed hybrid LSTM models that integrate technical indicators, sentiment, and macroeconomic data. These models provide a more comprehensive view of stock price drivers. Combining diverse data improves robustness and accuracy. The approach reflects real market conditions more realistically.
- [9] **Sirignano, J., Cont, R. (2015). *Universal Features of Price Formation in Financial Markets: Perspectives from Deep Learning. Quantitative Finance*, 19(9), 1449-1459.** They discussed the scalability issues of LSTM models in real-time stock forecasting. Most models focus on a few stocks and offline training, limiting practical use. Real-time systems require efficient retraining and fast inference. Bridging this gap between research and deployment remains a key challenge.
- [10] **Patel, J., Shah, S., Thakkar, P., Kotecha, K. (2015). *Predicting Stock Market Index Using Fusion of Machine Learning Techniques. Expert Systems with Applications*, 42(4), 2162-2172.** This study explored combining ML techniques to enhance stock index prediction. It emphasized the need to scale models for multiple stocks and real-time use. It also highlighted the importance of incorporating diverse, high-frequency data. Future models should unify attention, bidirectionality, and hybrid data in one efficient framework.

- [11] **Gareja Pradip, Chitrak Bari, J. Shiva Nandhini, "Stock market prediction using machine learning" International Journal of Advance Research and Development, Volume 3, Issue 10, 2014** This paper applied various machine learning algorithms like Linear Regression, Decision Trees, and SVM to predict stock prices. The study demonstrated improved accuracy over traditional models but emphasized the importance of selecting the right algorithm for specific datasets.
- [12] **K. Raza, "Prediction of Stock Market performance by using machine learning techniques", 2014.** Raza explored the application of ML techniques such as Decision Trees, Naive Bayes, and k-NN for predicting stock market performance. The study highlighted that ML models can capture non-linear relationships and improve prediction accuracy but are highly dependent on data quality.
- [13] **K. Hiba Sadia, Aditya Sharma, Adarrsh Paul, Sarmistha Padhi, Saurav Sanyal, "Stock Market Prediction Using Machine Learning Algorithms", International Journal of Engineering and Advanced Technology (IJEAT), Volume-8 Issue-4, 2014.** This work implemented machine learning algorithms like Random Forest and Gradient Boosting for stock market prediction. It showed that ensemble models perform better due to their ability to reduce overfitting and generalize well across different market conditions.
- [14] **Raut Sushrut Deepak, Shinde Isha Uday, Dr. D. Malathi, "MACHINE LEARNING APPROACH IN STOCK MARKET PREDICTION", 2011.** The authors reviewed various ML models and concluded that their predictive performance depends on proper data preprocessing and feature selection. They emphasized the growing relevance of AI in financial analytics and recommended hybrid approaches for improved results.
- [15] **Mehtab, S., Sen, J., A robust predictive model for stock price prediction using deep learning and natural language processing. In: Proceedings of the 7th International Conference on Business Analytics and Intelligence, 199, 2011.** This paper proposed a robust deep learning model that combines LSTM networks with natural language processing (NLP) techniques. It integrated financial news sentiment with price data, significantly enhancing prediction accuracy by capturing both quantitative and qualitative signals.

3. SOFTWARE REQUIREMENT ANALYSIS

3.1 Problem Statement

The stock market is inherently volatile and influenced by a multitude of complex, non-linear factors such as historical trends, investor sentiment, economic indicators, and global events. Traditional statistical and some machine learning models fail to capture the temporal dependencies and intricate patterns within stock price movements, leading to suboptimal forecasting accuracy. There is a need for a more advanced, scalable, and accurate prediction model that can effectively learn from sequential data and adapt to dynamic market conditions. This project aims to address this challenge by developing a stock price prediction system using advanced LSTM-based deep learning techniques—including Bidirectional LSTM, Stacked LSTM, and Attention Mechanisms—to enhance the prediction performance and interpretability across multiple stocks in real time.

3.2 Modules and their Functionalities

The stock price prediction system is composed of several interconnected modules, each responsible for a specific set of tasks to ensure accurate and efficient forecasting. The Data Collection Module is the foundation, responsible for fetching historical and real-time stock market data using APIs like yfinance. This data is then processed by the Data Preprocessing Module, which cleans and formats it, handles missing values, applies feature scaling (such as MinMaxScaler), and generates sequential input for the LSTM models.

To enhance prediction quality, a Feature Engineering Module may be included. This module introduces additional features such as technical indicators (e.g., moving averages, RSI) or external data like financial news sentiment, enriching the dataset with more context for learning. The Model Building Module follows, where different deep learning models are constructed using architectures like basic LSTM, Bidirectional LSTM, Stacked LSTM, and LSTM with Attention Mechanisms. These models are compiled using appropriate loss functions and optimizers tailored for time series forecasting.

Once built, the models are trained and assessed in the Model Training and Evaluation Module. This component ensures the model learns effectively from the data while avoiding overfitting through techniques like dropout and early stopping. Performance is measured using metrics such as RMSE

(Root Mean Squared Error), MAE (Mean Absolute Error), and MAPE (Mean Absolute Percentage Error). The trained models then pass into the Prediction Module, which performs both single-step and multi-step forecasts, generating future stock prices based on the input sequences.

To provide users with intuitive insights, the Visualization Module plots actual vs. predicted stock prices using visualization tools like Matplotlib or Plotly. If the model uses attention mechanisms, it can also visualize the attention weights to explain the model's focus on specific time steps. For user accessibility and interaction, an optional Web Application or User Interface Module can be developed using Streamlit or Flask. This module allows users to input stock symbols, set prediction parameters, and visualize results in real time through a responsive dashboard.

3.3 Non – Functional Requirements:

The stock price prediction system is designed not only for accurate forecasting but also to meet several key non-functional requirements that ensure its overall effectiveness and usability. Performance is a critical factor; the system must provide timely predictions with minimal latency, especially when working with real-time stock market data. Scalability is equally important—the architecture should be capable of handling large datasets and multiple stock symbols, as well as accommodating the integration of additional features like sentiment analysis or macroeconomic indicators in the future.

To ensure a positive user experience, the system should also prioritize usability. If implemented as a web application, the interface must be intuitive and user-friendly, allowing users to input stock symbols, set date ranges, and view forecasts and visualizations easily. Maintainability is another vital aspect; the code should be modular, well-documented, and follow standard practices to support future development and enhancements. Reliability is essential to prevent disruptions—robust error handling must be in place to manage issues like API failures or missing data gracefully. Finally, the system must maintain a high degree of accuracy by using well-tuned, state-of-the-art deep learning models, ensuring that the forecasts remain trustworthy for decision-making purposes.

3.4 Feasibility Study:

A feasibility study evaluates the practicality and effectiveness of implementing a scalable LSTM-based system for stock price forecasting. It considers technical, operational, economic, and legal aspects to ensure the project's success.

Technical Feasibility:

The use of LSTM networks is well-supported by modern deep learning frameworks such as TensorFlow and PyTorch. These tools enable efficient implementation of complex architectures, including attention mechanisms and multi-layer models. The availability of high-performance computing resources, including GPUs and cloud services, ensures the scalability of the model for handling large datasets and real-time predictions.

Operational Feasibility:

The proposed system can be seamlessly integrated into existing financial analysis workflows. It is user-friendly and designed to provide interpretable forecasts, which makes it suitable for traders, investors, and analysts.

Moreover, the modular structure ensures easy maintenance, updates, and extension of features.

Economic Feasibility:

Initial costs may include cloud computing subscriptions, data access fees, and development time. However, the potential return on investment (ROI) is high due to the system's capability to improve decision-making, reduce financial risk, and offer competitive advantages. Open-source tools and datasets can further reduce costs.

Legal and Ethical Feasibility:

The system will adhere to financial data usage policies and privacy regulations. It does not involve any sensitive user data and operates on publicly available market information, ensuring compliance with legal standards.

4. SOFTWARE AND HARDWARE REQUIREMENTS

4.1. SOFTWARE REQUIREMENTS

The software requirements define the tools, libraries, and environments necessary to develop and run the scalable time series forecasting system effectively. The proposed system leverages deep learning and data processing capabilities to predict stock prices. Key software components include:

- Python 3.8 or later (with IDLE or Jupyter Notebook)
- PyCharm or VS Code as the development environment
- TensorFlow or PyTorch for implementing LSTM models
- NumPy, Pandas, Scikit-learn for data handling and preprocessing
- Matplotlib, Seaborn for data visualization
- yFinance or Alpha Vantage API for stock data retrieval

4.2 HARDWARE REQUIREMENTS

The hardware configuration should support deep learning model training and large-scale time series processing. Minimum requirements include:

- Operating System: Windows 10/11 or Linux (Ubuntu preferred)
- Processor: Intel i5 or higher
- RAM: 8GB minimum (16GB recommended)
- Storage: 500GB HDD or SSD (SSD preferred)
- GPU (optional but recommended): NVIDIA CUDA-enabled GPU for faster model training

5. SYSTEM ARCHITECTURE

5.1 SYSTEM ARCHITECTURE




Fig. 5.1 System Architecture

5.2 Data Flow Diagram

The DFD is also called as bubble chart. It is a simple graphical formalism that can be used to represent a system in terms of input data to the system, various processing carried out on this data, and the output data is generated by this system. The data flow diagram (DFD) is one of the most important modeling tools. It is used to model the system components. These components are the system process, the data used by the process, an external entity that interacts with the system and the information flows in the

system. DFD shows how the information moves through the system and how it is modified by a series of transformations. It is a graphical technique that depicts information flow and the transformations that are applied as data moves from input to output. DFD is also known as bubble chart. A DFD may be used to represent a system at any level of abstraction. DFD may be partitioned into levels that represent increasing information flow and functional detail.




Figure:5.2 Dataflow Diagram

5.3 UML Diagrams

UML is a standard language for specifying, visualizing, constructing, and documenting the artifacts of software systems. UML was created by the Object Management Group (OMG) and UML 1.0 specification draft was proposed to the OMG in January 1997. There are several types of UML diagrams and each one of them serves a different purpose regardless of whether it is being designed before the implementation or after (as part of documentation). UML has a direct relation with object-oriented analysis and design. After some standardization, UML has become an OMG standard. The two broadest categories that encompass all other types are:

- Behavioral UML diagram and
- Structural UML diagram.

As the name suggests, some UML diagrams try to analyses and depict the structure of a system or process, whereas other describe the behavior of the system, its actors, and its building components.

Goals: The Primary goals in the design of the UML are as follows:

- Provide users a ready-to-use, expressive visual modeling Language so that they can develop and exchange meaningful models.
- Provide a formal basis for understanding the modeling language.
- Encourage the growth of OO tools market.
- Integrate best practices.

The different types are as follows:

- Sequence diagram
- Use case Diagram
- Activity diagram
- Class diagram
- Collaboration diagram

SEQUENCE DIAGRAM

A sequence diagram simply depicts interaction between objects in a sequential order i.e., the order in which these interactions take place. We can also use the terms event diagrams or event scenarios to refer to a sequence diagram. Sequence diagrams describe how and in what order the objects in a system function. These diagrams are widely used by businessmen and software developers to document and understand requirements for new and existing systems.




Figure 5.3.1 Sequence Diagram

USE CASE DIAGRAM

A use case diagram at its simplest is a representation of a user's interaction with the system that shows the relationship between the user and the different use case in which the user is involved. A use case diagram is used to structure of the behavior thing in a model. The use cases are represented by either circles or ellipses.




Fig 5.3.2 Use case diagram

ACTIVITY DIAGRAM

Activity diagram is another important diagram in UML to describe the dynamic aspects of the system. Activity diagram is basically a flowchart to represent the flow from one activity to another activity. This flow can be sequential, branched, or concurrent. Activity diagrams deal with all type of flow control by using different elements such as fork, join, etc.





Fig 5.3.3 Activity diagram

CLASS DIAGRAM

In software engineering, a class diagram in the Unified Modelling Language (UML) is a type of static structure diagram that describes the structure of a system by showing the system's classes, their attributes, operations (or methods), and the relationships among the classes. It explains which class contains information.



6. IMPLEMENTATION AND CODING

6.1 SOURCE CODE

TATA MOTORS Stock price prediction

```
import pandas as pd
import numpy as np
from sklearn.preprocessing
import MinMaxScaler
from tensorflow.keras.models
import Sequential
from tensorflow.keras.layers
import LSTM, Dense
import matplotlib.pyplot as plt

# Load the dataset
df = pd.read_excel('/content/tata motors (1).xlsx')

# Clean column names
df.columns = df.columns.str.strip().str.lower()

# Use only the 'close' column
df = df[['close']]

# Normalize the close prices scaler = MinMaxScaler()
df['close'] = scaler.fit_transform(df[['close']])
```

```

# Function to create sequences
def create_sequences(data, seq_length):
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:i + seq_length])
        y.append(data[i + seq_length])
    return np.array(X), np.array(y)

# Set sequence length based on dataset size
seq_length = 5 # Adjusted for small dataset

# Create sequences
X, y = create_sequences(df['close'].values, seq_length)

# Reshape X for LSTM [samples, timesteps, features]
X = X.reshape((X.shape[0], X.shape[1], 1))

# Build LSTM model model = Sequential([
    LSTM(50, return_sequences=False, input_shape=(X.shape[1], 1)),
    Dense(1)
])

# Compile model
model.compile(optimizer='adam', loss='mean_squared_error')

# Train model
model.fit(X, y, epochs=20, batch_size=1, verbose=1)

# Predict next value
last_sequence = df['close'].values[-seq_length:]
last_sequence = last_sequence.reshape((1, seq_length, 1))

```

```

predicted_scaled = model.predict(last_sequence)
predicted_price = scaler.inverse_transform(predicted_scaled)

print("\n Predicted next day price:", predicted_price[0][0])

# Predict on the training set for plotting predictions = model.predict(X)
actual_prices = scaler.inverse_transform(y.reshape(-1, 1))
predicted_prices = scaler.inverse_transform(predictions)

# Plot actual vs predicted
plt.figure(figsize=(10, 5))
plt.plot(actual_prices, label='Actual Price')
plt.plot(predicted_prices, label='Predicted Price')
plt.title('Tata Motors Stock Price Prediction')
plt.xlabel('Time')
plt.ylabel('Price')
plt.legend()
plt.grid(True)
plt.show()

```

Basic LSTM stock price prediction

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing
import MinMaxScaler
from sklearn.metrics
import mean_squared_error, mean_absolute_error

```

```

from keras.models import Sequential
from keras.layers import LSTM, Dropout, Dense
import plotly.graph_objects as go

# Load and prepare data
df = pd.read_csv('/content/yahoo_stock.csv')

# Use actual file
df['Date'] = pd.to_datetime(df['Date'])
df.set_index('Date', inplace=True)

# Feature scaling
scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(df['Close'].values.reshape(-1, 1))

# Create dataset with time steps
def create_dataset(data, time_step=60):
    X, y = [], []
    for i in range(len(data) - time_step - 1):
        X.append(data[i:(i + time_step), 0])
        y.append(data[i + time_step, 0])
    return np.array(X), np.array(y)

time_step = 60
train_size = int(len(scaled_data) * 0.8)
train_data = scaled_data[:train_size]
test_data = scaled_data[train_size:]

X_train, y_train = create_dataset(train_data, time_step)
X_test, y_test = create_dataset(test_data, time_step)

```

```

# Reshape for LSTM input: [samples, time steps, features]
X_train = X_train.reshape(-1, time_step, 1)
X_test = X_test.reshape(-1, time_step, 1)

# Build Basic LSTM model
model = Sequential()
model.add(LSTM(50, return_sequences=True, input_shape=(time_step, 1)))
model.add(Dropout(0.2))
model.add(LSTM(50, return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(50))
model.add(Dropout(0.2))
model.add(Dense(25))
model.add(Dense(1))

# Compile and train
model.compile(optimizer='adam', loss='mean_squared_error')
history = model.fit(X_train, y_train, batch_size=1, epochs=10)

# Predictions
train_predict = model.predict(X_train)
test_predict = model.predict(X_test)

# Inverse scaling
train_predict = scaler.inverse_transform(train_predict)
test_predict = scaler.inverse_transform(test_predict)
y_train_actual = scaler.inverse_transform([y_train])
y_test_actual = scaler.inverse_transform([y_test])

# Evaluation
train_rmse = np.sqrt(mean_squared_error(y_train_actual[0], train_predict[:, 0]))

```

```

train_mae = mean_absolute_error(y_train_actual[0], train_predict[:, 0])
test_rmse = np.sqrt(mean_squared_error(y_test_actual[0], test_predict[:, 0]))
test_mae = mean_absolute_error(y_test_actual[0], test_predict[:, 0])

print(f"Train RMSE: {train_rmse:.4f}, MAE: {train_mae:.4f}")
print(f"Test RMSE: {test_rmse:.4f}, MAE: {test_mae:.4f}")

# Plot
train_plot = np.empty_like(scaled_data)
train_plot[:, :] = np.nan
train_plot[time_step:len(train_predict)+time_step, :] = train_predict

test_plot = np.empty_like(scaled_data)
test_plot[:, :] = np.nan
test_plot[len(train_predict)+(time_step*2)+1:len(scaled_data)-1, :] = test_predict

fig = go.Figure()
fig.add_trace(go.Scatter(x=df.index, y=df['Close'], name='Actual', line=dict(color='blue')))
fig.add_trace(go.Scatter(x=df.index, y=train_plot[:, 0], name='Train Predict',
line=dict(color='green')))
fig.add_trace(go.Scatter(x=df.index, y=test_plot[:, 0], name='Test Predict', line=dict(color='red')))
fig.update_layout(title='Basic LSTM Stock Prediction', xaxis_title='Date', yaxis_title='Stock Price',
template='plotly_dark')
fig.show()

```

Bidirectional LSTM stock price prediction

```

import numpy as np
import pandas as pd
from sklearn.preprocessing

```

```

import MinMaxScaler
from sklearn.metrics
import mean_squared_error, mean_absolute_error
from keras.models
import Sequential
from keras.layers
import LSTM, Dense, Dropout, Bidirectional
import plotly.graph_objects as go

# Load and prepare data
df = pd.read_csv('/content/yahoo_stock.csv')
# Replace with your actual CSV
df['Date'] = pd.to_datetime(df['Date']) df.set_index('Date', inplace=True)

# Feature scaling
scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(df['Close'].values.reshape(-1, 1))

# Create dataset function
def create_dataset(data, time_step=60):
    X, y = [], []
    for i in range(len(data) - time_step - 1):
        X.append(data[i:(i + time_step), 0])
        y.append(data[i + time_step, 0])
    return np.array(X), np.array(y)

# Define time steps and split data
time_step = 60
train_size = int(len(scaled_data) * 0.8)
train_data = scaled_data[:train_size]
test_data = scaled_data[train_size:]

```

```

X_train, y_train = create_dataset(train_data, time_step)
X_test, y_test = create_dataset(test_data, time_step)

# Reshape to [samples, time steps, features]
X_train = X_train.reshape(-1, time_step, 1)
X_test = X_test.reshape(-1, time_step, 1)

# Build Bidirectional LSTM
model = Sequential()
model.add(Bidirectional(LSTM(50, return_sequences=True), input_shape=(time_step, 1)))
model.add(Dropout(0.2))
model.add(Bidirectional(LSTM(50, return_sequences=True))) model.add(Dropout(0.2))
model.add(Bidirectional(LSTM(50))) model.add(Dropout(0.2)) model.add(Dense(25))
model.add(Dense(1))

# Compile and train
model.compile(optimizer='adam', loss='mean_squared_error')
history = model.fit(X_train, y_train, batch_size=1, epochs=10)

# Predictions
train_predict = model.predict(X_train)
test_predict = model.predict(X_test)

# Inverse scaling
train_predict = scaler.inverse_transform(train_predict) test_predict =
scaler.inverse_transform(test_predict)
y_train_actual = scaler.inverse_transform([y_train]) y_test_actual =
scaler.inverse_transform([y_test])

```

```

# Evaluation

train_rmse = np.sqrt(mean_squared_error(y_train_actual[0], train_predict[:, 0]))
train_mae = mean_absolute_error(y_train_actual[0], train_predict[:, 0])
test_rmse = np.sqrt(mean_squared_error(y_test_actual[0], test_predict[:, 0]))
test_mae = mean_absolute_error(y_test_actual[0], test_predict[:, 0])

print(f'Train RMSE: {train_rmse:.4f}, MAE: {train_mae:.4f}')
print(f'Test RMSE: {test_rmse:.4f}, MAE: {test_mae:.4f}')

# Plotting results

train_plot = np.empty_like(scaled_data)
train_plot[:, :] = np.nan
train_plot[time_step:len(train_predict)+time_step, :] = train_predict

test_plot = np.empty_like(scaled_data)
test_plot[:, :] = np.nan
test_plot[len(train_predict)+(time_step*2)+1:len(scaled_data)-1, :] = test_predict

fig = go.Figure()
fig.add_trace(go.Scatter(x=df.index, y=df['Close'], name='Actual Price', line=dict(color='blue')))
fig.add_trace(go.Scatter(x=df.index, y=train_plot[:, 0], name='Train Predict',
line=dict(color='green')))
fig.add_trace(go.Scatter(x=df.index, y=test_plot[:, 0], name='Test Predict', line=dict(color='red')))
fig.update_layout(title='Bidirectional LSTM Stock Price Prediction', xaxis_title='Date',
yaxis_title='Stock Price', template='plotly_dark')
fig.show()

```

Stacked Bidirectional LSTM stock price prediction

```

import numpy as np
import pandas as pd
from sklearn.preprocessing
import MinMaxScaler
from sklearn.metrics
import mean_squared_error, mean_absolute_error
from tensorflow.keras.models
import Sequential
from tensorflow.keras.layers
import Dense, Dropout, Bidirectional, LSTM
import plotly.graph_objects as go

# Assuming df is your DataFrame with 'Close' prices and datetime index

# 1. Feature scaling
scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(df['Close'].values.reshape(-1, 1))

# 2. Prepare the dataset for LSTM
def create_dataset(data, time_step=1):
    X, y = [], []
    for i in range(len(data) - time_step - 1):
        X.append(data[i:(i + time_step), 0])
        y.append(data[i + time_step, 0])
    return np.array(X), np.array(y)

time_step = 60

# 3. Split data into train and test
train_size = int(len(scaled_data) * 0.8)

```

```

test_size = len(scaled_data) - train_size
train_data, test_data = scaled_data[0:train_size], scaled_data[train_size:]

# 4. Create datasets
X_train, y_train = create_dataset(train_data, time_step)
X_test, y_test = create_dataset(test_data, time_step)

# 5. Reshape input for LSTM [samples, time steps, features]
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1], 1)
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], 1)

# 6. Build Stacked Bidirectional LSTM Model
model = Sequential()
model.add(Bidirectional(LSTM(50, return_sequences=True), input_shape=(time_step, 1)))
model.add(Dropout(0.2))
model.add(Bidirectional(LSTM(50, return_sequences=True)))
model.add(Dropout(0.2))
model.add(Bidirectional(LSTM(50)))
model.add(Dropout(0.2))
model.add(Dense(25))
model.add(Dense(1))

# 7. Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

# 8. Train the model
history = model.fit(X_train, y_train, epochs=10, batch_size=1, verbose=1)

# 9. Predictions

```

```

train_predict = model.predict(X_train)
test_predict = model.predict(X_test)
# 10. Inverse transform to original scale
train_predict = scaler.inverse_transform(train_predict)
test_predict = scaler.inverse_transform(test_predict)
y_train_orig = scaler.inverse_transform(y_train.reshape(-1, 1))
y_test_orig = scaler.inverse_transform(y_test.reshape(-1, 1))

# 11. Calculate RMSE and MAE
train_rmse = np.sqrt(mean_squared_error(y_train_orig, train_predict))
train_mae = mean_absolute_error(y_train_orig, train_predict)
test_rmse = np.sqrt(mean_squared_error(y_test_orig, test_predict))
test_mae = mean_absolute_error(y_test_orig, test_predict)

print(f'Train RMSE: {train_rmse}, Train MAE: {train_mae}') print(f'Test RMSE: {test_rmse}, Test
MAE: {test_mae}')

# 12. Prepare data for plotting train_plot = np.empty_like(scaled_data)
train_plot[:, :] = np.nan
train_plot[time_step:len(train_predict) + time_step, 0] = train_predict[:, 0]

test_plot = np.empty_like(scaled_data) test_plot[:, :] = np.nan
test_plot[len(train_predict) + (time_step * 2) + 1:len(scaled_data) - 1, 0] = test_predict[:, 0]

# 13. Plot the results
fig = go.Figure()
fig.add_trace(go.Scatter(x=df.index, y=df['Close'], mode='lines', name='Actual Price',
line=dict(color='blue'))))
fig.add_trace(go.Scatter(x=df.index, y=train_plot[:, 0], mode='lines', name='Train Predict',
line=dict(color='green'))))

```

```
fig.add_trace(go.Scatter(x=df.index, y=test_plot[:, 0], mode='lines', name='Test Predict',  
line=dict(color='red')))
```

```
fig.update_layout(title='Stock Price Prediction using Stacked Bidirectional LSTM',  
xaxis_title='Date', yaxis_title='Stock Price', template='plotly_dark')
```

```
fig.show()
```

Attention based LSTM stock price prediction

```
import numpy as np  
import pandas as pd  
from sklearn.preprocessing  
import MinMaxScaler  
from sklearn.metrics  
import mean_squared_error, mean_absolute_error  
from tensorflow.keras.models  
import Model  
from tensorflow.keras.layers  
import Input, LSTM, Dense, Dropout, Layer, Multiply, Permute, RepeatVector, Lambda  
from tensorflow.keras  
import backend as K  
import plotly.graph_objects as go  
  
# 1. Data scaling and preparation (same as before)  
scaler = MinMaxScaler(feature_range=(0, 1))  
scaled_data = scaler.fit_transform(df['Close'].values.reshape(-1, 1))  
  
def create_dataset(data, time_step=1):  
    X, y = [], []
```

```

    for i in range(len(data) - time_step - 1):
        X.append(data[i:(i + time_step), 0])
        y.append(data[i + time_step, 0])
    return np.array(X), np.array(y)

time_step = 60
train_size = int(len(scaled_data) * 0.8)
train_data, test_data = scaled_data[:train_size], scaled_data[train_size:]

X_train, y_train = create_dataset(train_data, time_step)
X_test, y_test = create_dataset(test_data, time_step)

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1], 1)
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], 1)

# 2. Custom Attention Layer
class Attention(Layer):
    def __init__(self, **kwargs):
        super(Attention, self).__init__(**kwargs)

    def build(self, input_shape):
        self.W = self.add_weight(name='attention_weight', shape=(input_shape[-1], 1),
                                initializer='random_normal', trainable=True)
        self.b = self.add_weight(name='attention_bias', shape=(input_shape[1], 1),
                                initializer='zeros', trainable=True)
        super(Attention, self).build(input_shape)

    def call(self, x):
        # Alignment
        scores. Shape: (batch_size, time_steps, 1)

```

```

        e = K.tanh(K.dot(x, self.W) + self.b)
        # Attention
        weights. Shape: (batch_size, time_steps, 1)
        a = K.softmax(e, axis=1)
        # Weighted sum of inputs
        output = x * a
        return K.sum(output, axis=1)

# 3. Build model with attention
inputs = Input(shape=(time_step, 1))
lstm_out = LSTM(50, return_sequences=True)(inputs) drop1 = Dropout(0.2)(lstm_out)
lstm_out2 = LSTM(50, return_sequences=True)(drop1) drop2 = Dropout(0.2)(lstm_out2)
lstm_out3 = LSTM(50, return_sequences=True)(drop2) drop3 = Dropout(0.2)(lstm_out3)

attention_out = Attention()(drop3)

# Apply attention here

dense1 = Dense(25, activation='relu')(attention_out) output = Dense(1)(dense1)

model = Model(inputs=inputs, outputs=output)

model.compile(optimizer='adam', loss='mean_squared_error')

# 4. Train
history = model.fit(X_train, y_train, epochs=10, batch_size=1, verbose=1)

# 5. Predict and inverse scale
train_predict = model.predict(X_train)
test_predict = model.predict(X_test)

```

```

train_predict = scaler.inverse_transform(train_predict)
test_predict = scaler.inverse_transform(test_predict)
y_train_orig = scaler.inverse_transform(y_train.reshape(-1, 1))
y_test_orig = scaler.inverse_transform(y_test.reshape(-1, 1))

# 6. Calculate metrics
train_rmse = np.sqrt(mean_squared_error(y_train_orig, train_predict))
train_mae = mean_absolute_error(y_train_orig, train_predict)
test_rmse = np.sqrt(mean_squared_error(y_test_orig, test_predict))
test_mae = mean_absolute_error(y_test_orig, test_predict)

print(f'Train RMSE: {train_rmse}, Train MAE: {train_mae}')
print(f'Test RMSE: {test_rmse}, Test MAE: {test_mae}')

# 7. Prepare plot arrays
train_plot = np.empty_like(scaled_data)
train_plot[:, :] = np.nan
train_plot[time_step:len(train_predict) + time_step, 0] = train_predict[:, 0]

test_plot = np.empty_like(scaled_data)
test_plot[:, :] = np.nan
test_plot[len(train_predict) + (time_step * 2) + 1:len(scaled_data) - 1, 0] = test_predict[:, 0]

# 8. Plot results
fig = go.Figure()
fig.add_trace(go.Scatter(x=df.index, y=df['Close'], mode='lines', name='Actual Price',
line=dict(color='blue'))))

```

```

fig.add_trace(go.Scatter(x=df.index, y=train_plot[:, 0], mode='lines', name='Train Predict',
line=dict(color='green'))))
fig.add_trace(go.Scatter(x=df.index, y=test_plot[:, 0], mode='lines', name='Test Predict',
line=dict(color='red'))))

fig.update_layout(title='Stock Price Prediction using Attention-based LSTM',
axis_title='Date', yaxis_title='Stock Price', template='plotly_dark')
fig.show()

```

5 days FORECAST

```

import yfinance as yf
import numpy as np
import pandas as pd
from sklearn.preprocessing
import MinMaxScaler
from tensorflow.keras.models
import Sequential
from tensorflow.keras.layers
import Dense, LSTM, Dropout, Bidirectional
import matplotlib.pyplot as plt

# --- Hyperparameters ---
time_step = 60 ticker = 'AAPL'
forecast_days = 5

# --- Fetch real-time data ---
def fetch_stock_data(ticker='AAPL', period='180d', interval='1d'):
    df = yf.download(ticker, period=period, interval=interval, auto_adjust=True)
    df = df[['Close']]    df.dropna(inplace=True)

```

```

return df

# --- Create dataset ---
def create_dataset(data, time_step=1):
    X, y = [], []
    for i in range(len(data) - time_step - 1):
        X.append(data[i:(i + time_step), 0])
        y.append(data[i + time_step, 0])
    return np.array(X), np.array(y)

# --- Forecast future days ---
def forecast(model, last_seq, scaler, n=5):
    temp_input = list(scaler.transform(last_seq.reshape(-1, 1)))
    future_preds = []
    for _ in range(n):
        input_seq = np.array(temp_input[-60:]).reshape(1, 60, 1)
        pred = model.predict(input_seq, verbose=0)[0][0]
        future_preds.append(pred)
        temp_input.append([pred])
    return scaler.inverse_transform(np.array(future_preds).reshape(-1, 1))

# --- Main processing ---
df = fetch_stock_data(ticker)
data = df['Close'].values.reshape(-1, 1)
scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(data)

train_size = int(len(scaled_data) * 0.8)
train_data = scaled_data[:train_size]
test_data = scaled_data[train_size:]

```

```

X_train, y_train = create_dataset(train_data, time_step)
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1], 1)

# Only proceed if enough test data is available
if len(test_data) > time_step:
    X_test, y_test = create_dataset(test_data, time_step)
    X_test = X_test.reshape(X_test.shape[0], X_test.shape[1], 1)
else:
    print(" Not enough test data for the given time_step.")
    X_test, y_test = None, None

# --- Build Bidirectional LSTM model ---
model = Sequential()
model.add(Bidirectional(LSTM(50, return_sequences=True), input_shape=(time_step, 1)))
model.add(Dropout(0.2))
model.add(Bidirectional(LSTM(50, return_sequences=False))) model.add(Dropout(0.2))
model.add(Dense(1))

model.compile(optimizer='adam', loss='mean_squared_error')
model.fit(X_train, y_train, epochs=5, batch_size=32, verbose=1)

train_predict = model.predict(X_train, verbose=0)
train_predict = scaler.inverse_transform(train_predict)

# --- Evaluate and plot if test data is available ---
if X_test is not None:
    test_predict = model.predict(X_test, verbose=0)
    test_predict = scaler.inverse_transform(test_predict)
    y_test_actual = scaler.inverse_transform(y_test.reshape(-1, 1))

```

```

plt.figure(figsize=(10, 4))
plt.plot(df.index[-len(test_predict):], y_test_actual, label='Actual')
plt.plot(df.index[-len(test_predict):], test_predict, label='Predicted')
plt.title("Actual vs Predicted Prices (Test Set)")
plt.xlabel("Date")
plt.ylabel("Price")
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
else:
    print(" Skipping test set evaluation due to insufficient test data.")

# --- Forecasting ---
last_60_days = data[-60:]
future_forecast = forecast(model, last_60_days, scaler, forecast_days)

# --- Future forecast plot ---
future_dates = pd.date_range(df.index[-1], periods=forecast_days + 1, freq='D')[1:]
plt.figure(figsize=(8, 3))
plt.plot(future_dates, future_forecast, marker='o', linestyle='--', color='orange')
plt.title("Next 5 Days Forecast")
plt.ylabel("Price")
plt.xlabel("Date")
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Future forecast table ---

```

```
forecast_df = pd.DataFrame({"Date": future_dates.strftime("%Y-%m-%d"), "Predicted Price":  
future_forecast.flatten()})  
print(forecast_df)
```

6.2 IMPLEMENTATION

6.2.1 PYTHON

Python is a general-purpose interpreted, interactive, object-oriented, and high-level programming language. An Interpreted Language, Python has a design philosophy that emphasizes code readability (notably using whitespace indentation to delimit code blocks rather than curly brackets or keywords), and a syntax that allows programmers to express concepts in fewer lines of code than might be used in languages such as C++ or java. It provides constructs that enable clear programming on both small and large scales. Python interpreters are available for many operating systems.

Python, the reference implementation of Python, is open-source software and has a community-based development model, as do nearly all of its variant implementations. C Python is managed by the non - profit Python Software Foundation. Python features a dynamic type system and automatic memory management. It supports multiple programming paradigms, including object-oriented, imperative, functional and procedural, and has a large and comprehensive standard library. Python language is being used by almost all tech-giant companies like – Google, Amazon, Facebook, Instagram, Dropbox, Uber etc.

The biggest strength of Python is huge collection of standard libraries which can be used for the following –

- Speech Recognition
- GUI Applications (like Kivy, Tkinter, PyQt etc.)
- Machine Learning
- Test frameworks

- Multimedia

6.2.2 MACHINE LEARNING

Machine learning models for speech recognition require a large amount of data for training. This data includes audio samples of different users speaking various phrases and commands that the voice assistant should recognize. But in this project machine learning is playing a role indirectly through the use of two specific libraries: spaCy and the Google Cloud Speech Recognition API. Here, spaCy is used for natural language processing (NLP).

The `nlp` object is an instance of the spaCy language model, specifically the English language model (`en_core_web_sm`). When you process a text with `nlp`, it applies various machine learning models to tokenize, parse, and tag the input text. It identifies parts of speech, named entities, and other linguistic features.

The `recognize_google` method sends the recorded audio data to Google's servers, where machine learning models process the audio to transcribe it into text. This involves complex acoustic modeling and language modeling techniques. The Google Cloud Speech Recognition API uses machine learning models to handle various accents, intonations, and spoken language nuances, making it a powerful tool for speech recognition.

7. SYSTEM TESTING

The purpose of testing in scalable time series forecasting using advanced LSTM techniques is to evaluate the model's accuracy, robustness, and efficiency in predicting stock prices on unseen data,

ensuring it performs well across different market conditions and large datasets while optimizing hyperparameters for improved forecasting reliability and scalability.

7.1. Types Of Tests

Unit testing

Unit testing involves the design of test cases that validate that the internal program logic is functioning properly, and that program inputs produce valid outputs. All decision branches and internal code flow should be validated. It is the testing of individual software units of the application.

Unit tests ensure that each unique path of a business process performs accurately to the documented specifications and contains clearly defined inputs and expected results.

Inputs: contain all required fields (company Name, role, required Skills, and resume) with different scenarios to cover cases like missing company name, missing role, missing resume, and various skill checks.

Output: provide meaningful feedback according to the input scenario: confirmation of suitability or identification of missing fields or skills.

Functions: `checkResume()`, `document.getElementById()`, `split()`, `trim()`, `toLowerCase()`.

Integration testing

Integration tests are designed to test integrated software components to determine if they actually run as one program. Testing is event driven and is more concerned with the basic outcome of screens or fields. Integration tests demonstrate that although the components were individually satisfaction, as shown by successfully unit testing, the combination of components is correct and consistent.

Functional testing is centered on the following items:

Valid Input	: identified classes of valid input must be accepted
Invalid Input	: identified classes of invalid input must be rejected.
Functions	: identified functions must be exercised.
Output	: identified classes of application outputs must be exercised
Procedures	: interfacing systems or procedures must be invoked.

Organization and preparation of functional tests is focused on requirements, key functions, or special test cases. In addition, systematic coverage pertaining to identify Business process flows; data fields, predefined processes, and successive processes must be considered for testing. Before functional testing is complete, additional tests are identified and the effective value of current tests is determined.

System Test

System testing ensures that the entire integrated software system meets requirements. It tests a configuration to ensure known and predictable results. An example of system testing is the configuration-oriented system integration test. System testing is based on process descriptions and flows, emphasizing pre-driven process links and integration points.

White Box Testing

White Box Testing is a testing in which in which the software tester has knowledge of the inner workings, structure and language of the software, or at least its purpose. It is purpose. It is used to test areas that cannot be reached from a black box level.

Black Box Testing

Black Box Testing is testing the software without any knowledge of the inner workings, structure or language of the module being tested. Black box tests, as most other kinds of tests, must be written

from a definitive source document, such as specification or requirements document, such as specification or requirements document. It is a testing in which the software under test is treated, as a black box.

Unit Testing

Unit testing is usually conducted as part of a combined code and unit test phase of the software lifecycle, although it is not uncommon for coding and unit testing to be conducted as two distinct phases.

Test strategy and approach:

Field testing will be performed manually and functional tests will be written in detail.

Test objectives:

- All field entries must work properly.
- Pages must be activated from the identified link.
- The entry screen, messages and responses must not be delayed.

Integration Testing

Software integration testing is the incremental integration testing of two or more integrated software components on a single platform to produce failures caused by interface defects.

The task of the integration test is to check that components or software applications, e.g. components in a software system or – one step up – software applications at the company level – interact without error.

Test Results: All the test cases mentioned above passed successfully. No defects encountered.

Acceptance Testing

User Acceptance Testing is a critical phase of any project and requires significant participation by the end user. It also ensures that the system meets the functional requirements.

Test Results: All the test cases mentioned above passed successfully. No defects encountered.

7.2 Test cases

S.No	Test Case	Expected Result	Result	Remarks (If Fails)
1	Data Preprocessing and Feature Engineering	Cleaned, normalized, and appropriately scaled time series data with relevant features extracted (e.g., technical indicators, volume).	Pass	Data still contains outliers or missing values. Scaling method not optimal for LSTM.
2	Vanilla/Stacked LSTM Model Training	LSTM model (Vanilla or Stacked) with optimized hyperparameters trained on historical stock data.	Pass	Model fails to converge or exhibits vanishing/exploding gradients. Hyperparameters not optimized.
3	Bidirectional LSTM Model Training	Bidirectional LSTM model with optimized hyperparameters trained on historical stock data.	Pass	Bidirectional LSTM fails to converge or exhibits vanishing/exploding gradients. Training time excessively long.
4	Time Series Forecasting Accuracy (Short-Term)	Accurate short-term stock price predictions (e.g., next 1–5 days) with low Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE) on test data.	Pass	Significant deviation between predicted and actual short-term prices. MAE/RMSE values are high.

Table 7.2 Test Cases

8. OUTPUT

```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/rnn/rnn.py:200: UserWarning: Do not pass an "input_shape"/"input_dim" argument to a layer. When using Sequential models, prefer using an "Input(shape)" object as the first layer in the model instead.
  super().__init__(**kwargs)
Epoch 1/20      2s 4ms/step - loss: 0.2763
19/19      0s 3ms/step - loss: 0.0359
Epoch 2/20      0s 4ms/step - loss: 0.0222
19/19      0s 4ms/step - loss: 0.0218
Epoch 3/20      0s 5ms/step - loss: 0.0260
19/19      0s 5ms/step - loss: 0.0188
Epoch 4/20      0s 3ms/step - loss: 0.0176
19/19      0s 4ms/step - loss: 0.0209
Epoch 5/20      0s 3ms/step - loss: 0.0163
19/19      0s 3ms/step - loss: 0.0190
Epoch 6/20      0s 4ms/step - loss: 0.0217
19/19      0s 4ms/step - loss: 0.0244
Epoch 7/20      0s 3ms/step - loss: 0.0168
19/19      0s 4ms/step - loss: 0.0188
Epoch 8/20      0s 4ms/step - loss: 0.0206
19/19      0s 3ms/step - loss: 0.0161
Epoch 9/20      0s 4ms/step - loss: 0.0184
19/19      0s 5ms/step - loss: 0.0220
Epoch 10/20     0s 4ms/step - loss: 0.0124
19/19      0s 4ms/step - loss: 0.0173
1/1      0s 181ms/step
● Predicted next day price: 1088.8759
1/1      0s 278ms/step
```




Figure 8.1: Out Screen-1

```

Epoch 1/10
1399/1399 ————— 34s 22ms/step - loss: 0.0068
Epoch 2/10
1399/1399 ————— 31s 22ms/step - loss: 0.0014
Epoch 3/10
1399/1399 ————— 41s 22ms/step - loss: 0.0015
Epoch 4/10
1399/1399 ————— 41s 22ms/step - loss: 0.0014
Epoch 5/10
1399/1399 ————— 41s 22ms/step - loss: 0.0013
Epoch 6/10
1399/1399 ————— 41s 22ms/step - loss: 0.0013
Epoch 7/10
1399/1399 ————— 30s 22ms/step - loss: 0.0010
Epoch 8/10
1399/1399 ————— 42s 22ms/step - loss: 0.0011
Epoch 9/10
1399/1399 ————— 42s 23ms/step - loss: 9.5197e-04
Epoch 10/10
1399/1399 ————— 39s 21ms/step - loss: 0.0011
44/44 ————— 1s 22ms/step
10/10 ————— 0s 14ms/step
Train RMSE: 31.0845, MAE: 25.3412
Test RMSE: 70.4499, MAE: 55.2073

```




Figure 8.2: Out Screen-2

```

1399/1399 ————— 81s 54ms/step - loss: 0.0058
Epoch 2/10
1399/1399 ————— 80s 53ms/step - loss: 0.0015
Epoch 3/10
1399/1399 ————— 81s 52ms/step - loss: 0.0011
Epoch 4/10
1399/1399 ————— 83s 53ms/step - loss: 9.9553e-04
Epoch 5/10
1399/1399 ————— 81s 52ms/step - loss: 0.0013
Epoch 6/10
1399/1399 ————— 84s 53ms/step - loss: 9.3522e-04
Epoch 7/10
1399/1399 ————— 81s 53ms/step - loss: 8.7618e-04
Epoch 8/10
1399/1399 ————— 73s 52ms/step - loss: 8.1436e-04
Epoch 9/10
1399/1399 ————— 74s 53ms/step - loss: 9.8488e-04
Epoch 10/10
1399/1399 ————— 80s 52ms/step - loss: 7.3186e-04
44/44 ————— 2s 34ms/step
10/10 ————— 0s 21ms/step
Train RMSE: 41.9873, MAE: 33.6193
Test RMSE: 181.2181, MAE: 160.4631

```



Figure 8.3: Out Screen-3

```

1399/1399 ————— 84s 56ms/step - loss: 0.0048
Epoch 2/10
1399/1399 ————— 78s 56ms/step - loss: 0.0012
Epoch 3/10
1399/1399 ————— 81s 55ms/step - loss: 0.0016
Epoch 4/10
1399/1399 ————— 83s 56ms/step - loss: 0.0012
Epoch 5/10
1399/1399 ————— 80s 55ms/step - loss: 0.0011
Epoch 6/10
1399/1399 ————— 84s 57ms/step - loss: 9.2188e-04
Epoch 7/10
1399/1399 ————— 83s 58ms/step - loss: 0.0011
Epoch 8/10
1399/1399 ————— 79s 55ms/step - loss: 7.3197e-04
Epoch 9/10
1399/1399 ————— 81s 55ms/step - loss: 6.7705e-04
Epoch 10/10
1399/1399 ————— 83s 56ms/step - loss: 9.0163e-04
44/44 ————— 2s 36ms/step
10/10 ————— 0s 13ms/step
Train RMSE: 34.112508837056826, Train MAE: 18.08540393310447
Test RMSE: 68.83650267911583, Test MAE: 55.427784166837995

```



Figure 8.4: Out Screen-4

```

Epoch 1/10
1399/1399 ————— 35s 23ms/step - loss: 0.0136
Epoch 2/10
1399/1399 ————— 32s 23ms/step - loss: 0.0025
Epoch 3/10
1399/1399 ————— 32s 23ms/step - loss: 0.0022
Epoch 4/10
1399/1399 ————— 32s 23ms/step - loss: 0.0017
Epoch 5/10
1399/1399 ————— 42s 23ms/step - loss: 0.0018
Epoch 6/10
1399/1399 ————— 41s 23ms/step - loss: 0.0016
Epoch 7/10
1399/1399 ————— 40s 23ms/step - loss: 0.0011
Epoch 8/10
1399/1399 ————— 42s 23ms/step - loss: 0.0011
Epoch 9/10
1399/1399 ————— 41s 24ms/step - loss: 0.0011
Epoch 10/10
1399/1399 ————— 32s 23ms/step - loss: 8.4331e-04
44/44 ————— 1s 21ms/step
10/10 ————— 0s 14ms/step
Train RMSE: 44.56228384084889, Train MAE: 34.920865248255314
Test RMSE: 172.21953952000513, Test MAE: 157.95737095883018

```




Figure 8.5: Out Screen-5



Figure 8.6: Out Screen-6

9. CONCLUSION

This project decisively illustrates the immense potential of advanced Long Short-Term Memory (LSTM) networks in transforming stock price forecasting and enabling real-time market insights. By meticulously designing and optimizing architectures, including both unidirectional and bidirectional LSTMs, the solution effectively captures complex temporal dependencies and long-term patterns within highly volatile financial time series data. This capability leads to significantly improved predictive accuracy for short and medium-term horizons, surpassing the limitations of traditional statistical models.

Crucially, the focus on "scalability" ensures that the developed techniques can efficiently process and learn from vast quantities of historical and incoming real-time financial data across a multitude of stocks. This efficiency is paramount for practical deployment in dynamic trading environments, where timely predictions are critical. The successful integration of real-time forecasting further underscores the model's low-latency inference capabilities, vital for supporting automated trading strategies and providing up-to-the-minute market intelligence. While acknowledging the inherent unpredictability of stock markets due to external factors and the efficient market hypothesis, this work lays a robust foundation for leveraging cutting-edge AI to enhance decision-making in the financial sector, paving the way for more sophisticated algorithmic trading, risk management, and investment analysis tools. Continued research incorporating diverse data streams like news sentiment and macroeconomic indicators, alongside exploring even more advanced deep learning architectures, will further refine these capabilities.

10. FUTURE ENHANCEMENTS

Building upon the robust foundation established by advanced LSTM techniques for scalable and real-time stock price forecasting, several key future enhancements can significantly augment the system's predictive power, adaptability, and practical utility.

Firstly, integrating a broader spectrum of alternative data sources is paramount. Beyond traditional historical prices and technical indicators, incorporating news sentiment analysis, social media trends, macroeconomic indicators (e.g., interest rates, GDP growth, inflation), company earnings reports, and even satellite imagery for specific industries could provide richer contextual information. Deep learning models, particularly those leveraging attention mechanisms, can then be trained to weigh the importance of these diverse data streams dynamically, leading to more nuanced and robust predictions that account for external market drivers.

Secondly, exploring more sophisticated deep learning architectures beyond LSTMs holds immense promise. Transformer networks, with their self-attention mechanisms, have demonstrated superior performance in capturing long-range dependencies and parallelizing computations, potentially offering advantages in handling extremely long historical sequences and improving training efficiency.

Thirdly, enhancing the real-time adaptive capabilities of the models is crucial. Implementing online learning or transfer learning mechanisms would allow the models to continuously update and adapt to sudden shifts in market dynamics, unforeseen events, or changes in data distribution without requiring a full retraining from scratch. This would involve techniques like incremental learning or federated learning, where models are periodically updated with new data streams or trained on local datasets before merging insights globally.

Finally, improving model interpretability and explainability will be vital for fostering trust and adoption, especially in regulated financial environments. While deep learning models are often black boxes, techniques such as SHAP (SHapley Additive exPlanations) values or LIME (Local Interpretable Model-agnostic Explanations) can provide insights into which features most influence a particular prediction. Developing methods to visualize attention weights in Transformer models could also help identify which historical periods or data streams are deemed most important for a given forecast, making the model's reasoning more transparent and actionable for human analysts.

11. REFERENCES

- [1] Wang, X., Han, Y., Jiang, Q., & Ren, P. (2020). Stock price prediction using LSTM and Bi-LSTM models. *International Journal of Computer Science and Engineering*, 17(2), 263-269.
- [2] Livieris, I. E., Pintelas, P. E., & Pintelas, P. E. (2020). A novel hybrid forecasting model for stock price prediction using a deep learning LSTM and ARMA models. *Computational Intelligence and Neuroscience*, 2020.
- [3] Chen, R. T., Lai, J. S., & Chen, J. M. (2019). Financial time series forecasting with recurrent neural networks. *Expert Systems with Applications*, 108, 143-154.
- [4] Kim, J. (2019). Stock price prediction using LSTM, GRU, and attention mechanism. *Applied Soft Computing*, 83, 105658.
- [5] Ishita Parmar, Navanshu Agarwal, Sheirsh Saxena, Ridam Arora, Shikhin Gupta, Himanshu Dhiman, Lokesh Chouhan Department of Computer Science and Engineering National Institute of Technology, Hamirpur – 177005, 2019
- [6] Pranav Bhat Electronics and Telecommunication Department, Maharashtra Institute of Technology, Pune. Savitribai Phule Pune University, 2018
- [7] Anurag Sinha Department of computer science, Student, Amity University Jharkhand Ranchi, Jharkhand (India), 2018
- [8] V Kranthi Sai Reddy Student, ECM, Sreenidhi Institute of Science and Technology, Hyderabad, India, 2017

- [9] Y. Liang and D. Niu and M. Ye and W.C. Hong, Short-term load forecasting based on wavelet transform and least squares support vector machine optimized by improved cuckoo search, *Energies*, 9(10)(2016), 827

- [10] G. Li and J. Shi, On comparing three artificial neural networks for wind speed forecasting, *Applied Energy*, 87(7)(2016), pp. 2313-2320.

- [11] C.J. Lu and J.Y. Wu, An efficient cmac neural network for stock index forecasting, *Expert Systems with Applications*, 38(12) (2015), pp. 15194-15201.

- [12] J. Lu and Y. Bai, Applications of grnn based on particle swarm algorithm forecasting stock prices, *Perception & Psychophysics*, 30(3)(2015), pp. 303-307.

- [13] J.D. Mao and X.J. Zhang and J. Li, Wind power forecasting based on the bp neural network, *Applied Mechanics & Materials*, 341-342(341-342)(2013), pp. 1303-1307.

- [14] K.R. Mller and A.J. Smola and G. R?tsch and B. Sch?lkopf and J. Kohlmorgen and V. Vapnik, Predicting time series with support vector machines, *Advances in Kernel Methods-Support Vector Learning*, 20(2)(2013), pp. 999-1004.

- [15] E.E. Osuna, Support vector machines: training and applications, A. I. Memo no. 1602, C. B. C. L. Paper, 144(9)(2011), pp. 1308-16.